

A Tidy Approach to Text Analysis in R

ART Forum 2018

Preamble

- The tidy approach is only *one* way to do text analysis in R.
- My goal is to get you doing text analysis as quickly as possible.
- This is the *last* two weeks of my Marketing Analytics class.
- I assume you have basic fluency in R and the tidyverse.
- We'll focus on word frequencies, visualizations, sentiment analysis, and topic modeling.
- The slides (and cheatsheet) are designed to serve as a reference.
- Slides and data at github.com/marcdotson/tidy-text-analysis.

Tokenizing
and Visualization

Sentiment Analysis

Topic Modeling

Text as Unstructured Data

- Text is **unstructured**.
 - Authors can express themselves freely.
 - The same idea can be expressed in many ways.
- Text is increasingly available and important in marketing applications (e.g., social media, product reviews).

Working with Text

- The basic approach when analyzing text is called **bag of words**, where each word is considered separately (i.e., without **syntax**).
 - Each unique word is called a **term**.
 - Each realization of a term is called a **token**.
 - A **document** is written by an author.
 - A collection of documents is called a **corpus**.
- By *counting* the tokens for each term, we produce **word frequencies** that we can visualize.
- We can use word frequencies to determine document **sentiment** and apply *unsupervised learning* techniques to find **topics**.

Tidy Text

- The tidyverse represents the state of the art for data wrangling and summarization in R and serves as the *foundation for a growing number of additional packages*.
- The `tidytext` package is an ideal example of this: it utilizes the tidyverse packages and adheres to the tidy data philosophy:
 1. Each observation (*token*) has its own row.
 2. Each variable has its own column.
 3. Each value has its own cell.

Text, Characters, and Strings

- Load the `tidyverse` and install and load `tidytext`.

```
text <- c(
  "So long and thanks for all the fish,",
  "So sad that it should come to this,",
  "We tried to warn you all but oh dear!"
)
```

- What data type and structure is `text`?
- Turn `text` into a data frame (i.e., tibble) and number the lines. (Hint: Use `data_frame()`.)
- Is this tidy?

Tokenize

- Use `unnest_tokens()` to **tokenize** the text (i.e., split it into individual words or tokens).

```
text_df %>%  
  unnest_tokens(word, text)
```

- This is a *powerful* function.
 - Each token (i.e., word) has its own row.
 - The line number is preserved.
 - Tokens are converted to lowercase.
 - Whitespace and punctuation are stripped.
- Now that the data is tidy, all tidyverse tools are applicable!

Down the Rabbit Hole

- Install and load `gutenbergr`, a package that provides easy access to [Project Gutenberg](#) (the `gutenberg_id` is in a book's URL).
- Let's look at Lewis Carroll's *Alice's Adventures in Wonderland*.

```
tidy_carroll <- gutenbergr_download(11) %>%  
  unnest_tokens(word, text)
```

- Now that we have tidy data, how do we compute the word frequencies (in descending order)?
- Does this seem right to you?

Remove Stop Words

- Commonly used words aren't very informative and are referred to as **stop words**.

```
stop_words
```

- This is just a data frame, and we know how to join data frames!

```
tidy_excited <- tidy_excited %>%  
  anti_join(stop_words)
```

- Why didn't we need to specify **by**?
- Produce the word frequencies (in descending order) again.

Visualize Word Frequencies

- Frequencies are counts, which says use bar plots.

```
tidy_carroll %>%  
  count(word) %>%  
  mutate(word = reorder(word, n)) %>%  
  ggplot(aes(x = word, y = n)) +  
  geom_col()
```

- What can we do to make this plot readable?

Word Clouds

- With the **size of words indicating frequency**, a word cloud might also be a helpful visualization.
- Install and load the `wordcloud` package.

```
tidy_carroll %>%  
  count(word) %>%  
  with(wordcloud(word, n, min.freq = 10))
```

- Note that the location of a word in the cloud is random.

Exercise

Visit [Project Gutenberg](#) and select a book to analyze as follows.

1. In consecutive lines of code, download (the `gutenberg_id` is in a book's URL) the text, tokenize, and remove stop words.
2. Visualize the word frequencies using a column plot and a word cloud.

Tokenizing
and Visualization

Sentiment Analysis

Topic Modeling

Web Scraping

- Let's scrape data from the web; install and load `rvest`.

```
text <- read_html(  
  "https://en.wikipedia.org/wiki/Columbus,_Ohio"  
) %>%  
  html_nodes("#content") %>%  
  html_text() %>%  
  str_split("\\\\n\\\\n\\\\n") %>%  
  unlist()
```

- The **node** and **regular expression** are also specific to this webpage.

Tokenize, Tidy, and Visualize

- Turn `text` into a data frame, create a *section* variable, tokenize, and remove stop words.

```
tidy_text <- data_frame(text) %>%  
  mutate(section = row_number()) %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words)
```

- Now visualize the word frequencies.
- Counts suggest meaning, but what is the emotional content?

Sentiment Dictionaries

- **Sentiment** is a reference to the emotional content of words.
- Like bag of words, the basic approach to sentiment analysis is to use a **sentiment dictionary** (i.e., lexicon).

```
get_sentiments("afinn")
```

- Look at the **bing** and **nrc** sentiment dictionaries.
- What are the ten sentiments in the **nrc** sentiment dictionary?

Sentiment Analysis

- A sentiment dictionary is just a data frame, and we know how to join data frames!

```
sentiment_nrc <- tidy_text %>%  
  inner_join(get_sentiments("nrc"))
```

- What sentiments are represented most frequently in our data?
- What words contribute to the “joy” sentiment in our data?
- Note that a sentiment dictionary is time and application-specific.

Changing Sentiment

- How does sentiment change over a lengthy document?

```
tidy_carroll <- gutenbergs_download(11) %>%  
  mutate(line = row_number()) %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words)
```

- Let's build a pipeline into the visualization.

```
tidy_carroll %>%  
  inner_join(get_sentiments("bing")) %>%  
  count(index = line %% 30, sentiment) %>%  
  spread(sentiment, n, fill = 0) %>%  
  mutate(sentiment = positive - negative) %>%  
  ggplot(aes(x = index, y = sentiment)) +  
  geom_col()
```

Exercise

Scrape a Wikipedia page or download from [Project Gutenberg](#).

1. In consecutive lines of code, create sections of the text, tokenize, and remove stop words.
2. Append a sentiment dictionary. What sentiments are most frequently represented?
3. Visualize how sentiment changes across the document.

Tokenizing
and Visualization

Sentiment Analysis

Topic Modeling

Word Frequencies

- Let's compare two of Lewis Carroll's books: *Alice's Adventures in Wonderland* and *Through the Looking-Glass*.

```
tidy_carroll <- gutenbergl_download(c(11, 12)) %>%  
  unnest_tokens(word, text) %>%  
  mutate(  
    book = factor(  
      gutenbergl_id,  
      labels = c(  
        "Alice's Adventures in Wonderland",  
        "Through the Looking-Glass"  
      )  
    )  
  ) %>%  
  count(book, word) %>%  
  arrange(desc(n))
```

Term Frequency-Inverse Document Frequency

- The **tf-idf** statistic weights the word frequencies for each document by how uncommon the word is within the corpus.

```
tidy_carroll %>%  
  bind_tf_idf(word, book, n)
```

- Overwrite `tidy_carroll` with this new data frame and arrange in descending order by `tf-idf`.
- What words are most unique (i.e., important) for each book?

Visualize tf-idf by Document

- Let's create a single visualization for our corpus.

```
tidy_carroll %>%  
  group_by(book) %>%  
  top_n(10, tf_idf) %>%  
  ungroup() %>%  
  mutate(word = reorder(word, tf_idf)) %>%  
  ggplot(aes(word, tf_idf, fill = book)) +  
  geom_col(show.legend = FALSE) +  
  facet_wrap(~ book, scales = "free") +  
  coord_flip()
```



Beyond Word Frequencies

- Word frequencies (and tf-idf) suggest overall meaning and sentiment analysis describes emotional intent, but we would like to uncover the different **topics** being written about.
 - This is especially true as the size of a corpus increases.
- **Topic models** uncover groups of words (i.e., topics) via *unsupervised learning*.
- The most common topic model is called **latent Dirichlet allocation** or LDA.

Clustering vs. Topic Modeling

- Clustering
 - Clusters are uncovered based on *distance*, which is continuous.
 - Every object is assigned to a single cluster.
 - Clusters are summarized based on averages.
- Topic Modeling
 - Topics are uncovered based on *word frequency*, which is discrete.
 - Every document is a **mixture** (i.e., partial member) of every topic.
 - Topics are summarized based on word probabilities.

Latent Dirichlet Allocation



Create a Document Term Matrix

- Install and load `topicmodels`.
- Import `Roomba 650 Amazon Reviews.csv`, create a `review id` variable, tokenize, remove stop words, and select the `review` and `word` variables, and assign to `roomba_650`.
- The input for a topic model is not tidy data, it's a **document term matrix**. Let's “cast” our tidy data into a DTM.

```
dtm_text <- roomba_650 %>%  
  count(review, word) %>%  
  cast_dtm(review, word, n)
```

Run a Topic Model

- Running a topic model is straightforward with a DTM.

```
lda_out <- dtm_text %>%  
  LDA(  
    k = 2,  
    method = "Gibbs",  
    control = list(seed = 42)  
  )
```

Topic Word Probabilities

- The most important output from a topic model are the topics themselves: the dictionary of words, sorted according to the **probability** the word is part of that topic.
- Let's "tidy" these probabilities (i.e., betas).

```
lda_topics <- lda_out %>%  
  tidy(matrix = "beta")
```

Visualize, Name, and Choose K

- Now we can visualize our topics.

```
lda_topics %>%  
  group_by(topic) %>%  
  top_n(15, beta) %>%  
  ungroup() %>%  
  mutate(term = reorder(term, beta)) %>%  
  ggplot(aes(term, beta, fill = as.factor(topic))) +  
  geom_col(show.legend = FALSE) +  
  facet_wrap(~ topic, scales = "free") +  
  coord_flip()
```

- How would you name these two topics?
- How can you be sure $k = 2$ is enough topics?

Exercise

Using the `roomba_650` data, do the following.

1. Run a number of topic models with varying k .
2. Visualize and compare topics across solutions. Which solution is best?

Epilogue

- We've covered the basics of a tidy approach to text analysis.
- You've (hopefully) been convinced of the benefit of building an analysis on the tidyverse.
- The focus has been on word frequencies, visualizations, sentiment analysis, and topic modeling.

What's Next?

- *Text Mining with R*
- *R for Data Science*
- DataCamp

Bonus: HTML Nodes

- When scraping a webpage, you need to identify the **HTML node** you want to extract using a **selector**.
- An easy one to use is **SelectorGadget**.
- Use a selector to identify the node we want to extract from the Wikipedia entry for **The Hitchhiker's Guide to the Galaxy**.
- Load `tidyverse`, `tidytext`, and `rvest`.

```
read_html("<URL>") %>%  
  html_nodes("<NODE>") %>%  
  html_text()
```

Bonus: Regular Expressions

- *Regular expressions* (or regexps) are used to **describe patterns in strings** and are thus helpful in splitting a string into sections, lines, chapters, and documents. Some basics:
 - "abc" matches abc
 - "." matches any character
 - "\\s" matches whitespace
 - "\\d" matches any digit
 - "(a|b)" matches a or b
 - "\\" escapes special behavior (e.g., "\\." matches .)
- Practice writing regular expressions with `str_view_all()`.

```
test <- "Text in 452 is a textbook end to Winter 2018!"  
str_view_all(test, "end")
```

- Write a regexp to match both uses of “text” in `test`.